Style Forecast: Technical Report

COMP3011 Coursework 1

March 2026

GitHub Repository
API Documentation
Presentation Slides

1 Overview

Style Forecast is an API-first outfit recommendation system designed for a practical everyday problem: people often have limited time when getting ready, incomplete awareness of weather conditions at a chosen time, and only partial knowledge of current style trends. The system bridges that gap by combining live weather context, magazine-derived trend signals, and lightweight user context into a recommendation that helps the user feel both comfortable and confident.

The central design objective was not simply to return an outfit, but to return one that is appropriate for climate, relevant to current fashion context, and explainable enough to earn user trust. This emphasis on explainability shaped both the system architecture and the recommendation logic, and it differentiates the project from many fashion tools that prioritise inspiration or browsing over defensible recommendation behaviour.

2 Methodology

2.1 Problem Framing and Design Approach

The project was approached as a decision-support system rather than a generic fashion catalogue. A recommendation is only useful if it reduces uncertainty at the point of decision-making, so the system needs to reason about context rather than preference alone. For that reason, the chosen methodology combines:

- **weather suitability**, to reduce the risk of physical discomfort,
- **trend relevance**, to reflect recent fashion signals,
- **bounded scoring**, to keep the recommendation stable, explainable, and technically defensible.

Alternative approaches were considered before selecting the final model. Pure collaborative filtering was not suitable because the dataset is small and suffers from cold-start problems. An end-to-end LLM ranking approach was also rejected because it would be harder to audit, less deterministic, and less reliable for a live assessment setting. GraphQL was considered as an alternative API style, but the assessed workflows favour stable resources, explicit status-code behaviour, and straightforward Swagger documentation. A hybrid weighted model therefore offered the best balance of interpretability, control, and practical usefulness.

Style Forecast is not presented as the first system to combine fashion and weather cues. Its originality lies in how it integrates forecast context, magazine-derived trend signals, and bounded recommendation logic into a single explainable pipeline that can be defended technically. In

practice, a user could try to recreate a similar decision process manually with a general-purpose AI assistant such as ChatGPT, but that would require repeatedly gathering weather data, checking current fashion coverage, restating context, and interpreting unstructured responses each day. Style Forecast streamlines that fragmented task into a focused, repeatable API workflow with explicit inputs, auditable scoring, and an interface tailored to this specific decision.

2.2 Architecture Overview

The system is structured in three layers, shown in Figure 1. This separation keeps the application maintainable and makes design choices easier to justify.

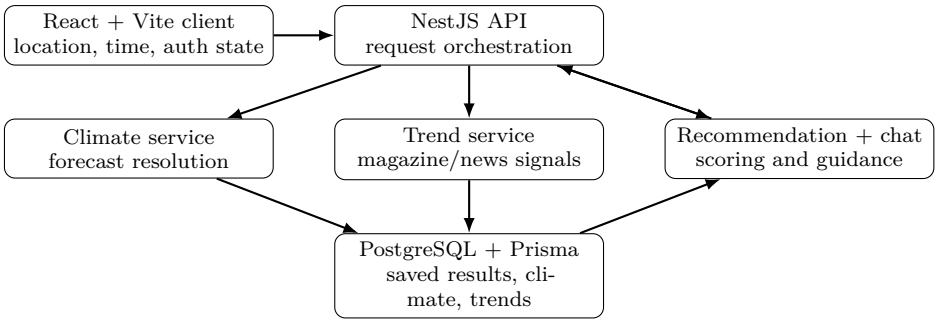

Figure 1: Simplified system flow for API orchestration, trend processing, persistence, and chat-assisted recommendation.

The client is intentionally thin: it collects user input, presents results, and manages authenticated state. The API is the coordination layer for validation, scoring, ownership rules, and external data access.

2.3 Technology Choices and Rationale

TypeScript was chosen across frontend and backend to reduce integration mismatch and improve refactor safety [1]. Using one typed language simplified DTO design, API contracts, and debugging.

NestJS was chosen for the backend because its modular structure fits a system with distinct domains such as authentication, climate retrieval, trend ingestion, persistence, and recommendation logic [2]. Keeping logic in services rather than controllers improved clarity and made it easier to adjust scoring logic without destabilising routing behaviour.

PostgreSQL with **Prisma** was chosen for persistence [4, 3]. An SQL database is appropriate because the project still has relational concerns: users, saved recommendation snapshots, climate records, and trend entries all benefit from structured querying and clear relationships. The schema also includes item and outfit resources used to satisfy the assessed CRUD requirement and support stored recommendation structures. Figure 2 summarises the core persistence model used by the app and backend services.

React + Vite was chosen for the client because the interaction model is stateful and iterative [6, 7]. React supports interactive recommendation flows and authenticated UI state, while Vite keeps iteration fast.

REST over GraphQL was also a deliberate choice. GraphQL would have been defensible for highly variable client-side queries, but this project has a relatively stable set of resources and workflows. REST made route semantics, status codes, authentication guards, and Swagger-based documentation more direct [5].

3 Implementation

3.1 Core Data, CRUD, and Resource Design

To satisfy the coursework requirement for database-backed CRUD, the backend exposes full RESTful CRUD for the `Item` model:

- `POST /api/items` creates an item,
- `GET /api/items` and `GET /api/items/:id` read user items,
- `PATCH /api/items/:id` updates an item,
- `DELETE /api/items/:id` deletes an item.

These operations are backed by PostgreSQL through Prisma and return conventional JSON status-code responses.

The `Item` model was the most appropriate candidate for full CRUD because it is a stable resource with clear lifecycle operations and conventional REST semantics. In the submitted app, however, CRUD is not the main user-facing story; the visible experience is recommendation-led, while these resources mainly support the assessed API surface and stored recommendation structures. The listings subsystem was also extended to support a fuller lifecycle: listings can be updated or deleted by their owner, and deleting a listing releases the linked item back to the `AVAILABLE` state instead of leaving it trapped in a reserved marketplace state.

3.2 Recommendation Pipeline

The recommendation service combines two primary signals. First, climate suitability measures whether a suggested look is likely to be comfortable for the chosen location and time. Second, trend relevance applies a bounded boost or penalty from recent magazine-derived fashion signals gathered from curated feeds. The recommendation endpoint remains publicly usable, while authenticated requests unlock saved recommendations, account settings, and more tailored chat assistance.

This design is especially useful in the target use case. A user getting ready quickly may know that the day is "cold" or "rainy", but not whether the weather will change by the time they leave. By consolidating these signals into one recommendation, the system reduces fragmented decision-making and supports comfort and self-presentation. The more original aspect of the design is that climate, trend, and preference signals are not merely displayed side by side; they are fused into one bounded scoring pipeline whose contribution can still be explained and defended.

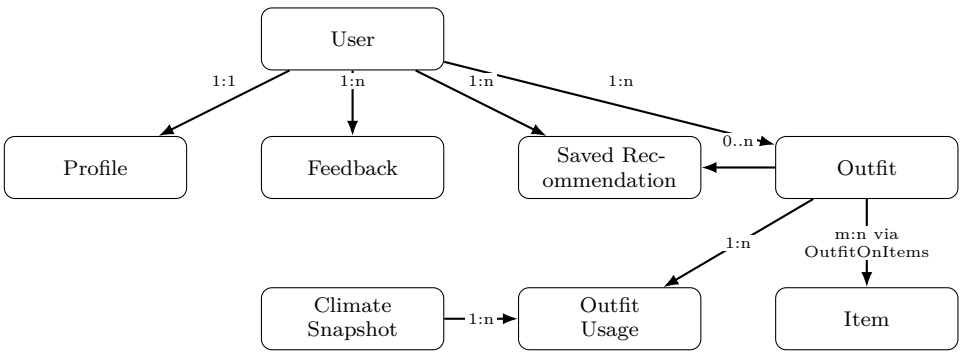

Figure 2: Core ER diagram for the database elements most relevant to authentication, recommendation context, saved results, and CRUD support. Marketplace and recycling tables are omitted for readability.

3.3 Data Flow Through the System

The operational flow begins when the user sends a recommendation request, optionally after logging in. The client collects a location and time, the API resolves weather context, and trend services contribute recent magazine-style signals. A saved account location can prefill the form, while blank settings intentionally do not force an autofill. After scoring candidate looks, the API returns a JSON response containing the selected result, and logged-in users can store that recommendation with its weather, place, time, stylist blurb, and a small deduped set of inspiration image URLs for later review. A separate chat-assistant flow can then turn the same context into more conversational styling advice. In the final frontend, that assistant is deliberately streamlined around preset activity selections rather than an unrestricted compose box, which reduces noisy prompts and keeps the guidance more consistent during demonstration.

3.4 Security, Error Handling, and Operational Choices

Security was addressed at multiple layers.

Authentication. The API uses JWT-based authentication with guarded routes. This provides a simple stateless security model for a coursework API and allows protected ownership checks without server-side session storage [9]. The account area also supports authenticated password change with current-password verification before a new bcrypt hash is stored. JWT secrets are treated as required environment configuration rather than silently falling back to hard-coded defaults.

Authorisation. Ownership checks were added to CRUD operations and saved recommendations so a user can only read, update, or delete their own records. This prevents horizontal privilege issues and ensures the API does not expose another user's private data.

Validation and error handling. DTO validation is enforced in NestJS using class-validator and a global validation pipe. Explicit `401`, `403`, `404`, and `409` responses make failure states predictable instead of ambiguous. This is particularly important for ownership-controlled endpoints, where incorrect status behaviour can hide security mistakes during testing.

Proxy policy. The image proxy used for Pinterest inspiration was tightened so it only accepts HTTPS requests to approved Pinterest image hosts, applies a timeout, and verifies that the upstream response is an image. This reduces the risk of turning a convenience proxy into a general server-side fetch endpoint.

Operational trade-off. JWT was selected because it is simple and stateless, but token revocation is weaker than in a full server-side session model. For this coursework-scale system, simplicity and ease of integration outweighed that limitation.

3.5 Documentation, Version Control, and Deployment Readiness

The project is documented in two complementary ways: Swagger-based API documentation for endpoint use, and this report for architectural rationale, testing, and limitations. The final submission was tightened so these artifacts describe the same recommendation-led system rather than contradictory scopes.

Version control was managed through GitHub using named branches and incremental commits for major features such as recommendation refinement, authentication hardening, chat-assistant integration, trend-driven interface changes, and documentation updates. This supported both experimentation and a cleaner final assessed surface.

The system is demonstrable through local execution. Backend and frontend builds run successfully, the API exposes JSON responses over HTTP, and a Swagger interface is available for interactive testing. These factors satisfy the coursework requirement that the system be runnable and inspectable rather than merely described.

4 Evaluation

4.1 Usefulness and Design Value

The system is useful because it reduces cognitive load during a time-constrained daily task. Users may know the weather in general terms but still miss forecast detail at the exact time they are going out, and may only have vague awareness of current fashion coverage. Style Forecast turns those fragmented signals into one actionable suggestion.

4.2 Testing and Validation

Evaluation was carried out in four ways:

1. backend and frontend production builds to confirm deployable integrity,
2. endpoint validation through Swagger and direct HTTP requests,
3. automated unit tests for recommendation and account logic,
4. end-to-end flow checks across registration, authenticated recommendation, saved recommendations, password change, and chat-assisted guidance.

The evaluation is qualitative rather than benchmark-driven: the main question is whether the recommender gives plausible, explainable, and secure recommendations under realistic usage conditions. In addition to unit tests, the assessed backend includes integration coverage for authenticated recommendation behaviour, saved recommendation flows, and chat-assisted guidance.

4.3 Challenges, Lessons, and Future Work

The most difficult implementation issue was timezone correctness. Weather relevance depends on when and where the user intends to wear an outfit, so naive datetime handling produces incorrect recommendations. This also affected account settings, because the interface needed to display an obvious user-facing timezone label while still storing a robust underlying zone identifier. A second challenge was external data reliability, so the system had to fail clearly rather than silently.

The main lesson from the project was that explainability matters as much as raw feature count. A smaller system with clear logic, secure defaults, and a strong end-to-end flow is more defensible than a larger but less coherent implementation. Architectural separation between controller, service, and persistence layers also proved worthwhile when CRUD, JWT refresh, and documentation changes had to be introduced without destabilising unrelated behaviour.

The current trend sentiment model is lexical and can miss nuance. Personalisation is heuristic rather than learned from large-scale data, and automated coverage remains narrower than production standards. Future work would therefore focus on broader automated testing, scheduled trend ingestion, constrained learning-to-rank, and richer experimentation with explanation-preserving ranking methods rather than opaque generation-heavy recommenders.

5 Conclusion

Style Forecast was designed to bridge a genuine everyday knowledge gap: what to wear when time is limited and contextual information is fragmented. The final architecture reflects deliberate choices about maintainability, security, explainability, documentation quality, and practical usefulness. The project therefore demonstrates not only a working API, but a defensible design process backed by version control, testing, and structured technical reasoning.

Appendix A Generative AI Declaration and Conversation-Log Summary

Generative AI tools were used selectively for ideation, high-level alternative exploration, implementation checking, report restructuring, and presentation refinement. AI was not treated as an authoritative source; final architecture, implementation, testing, and submission decisions were reviewed and verified manually.

Main GenAI Uses

- exploring alternative high-level project directions and architectures, including climate-aware fashion intelligence, multi-objective scoring, trend integration, GraphQL, collaborative filtering, and LLM-based ranking, before selecting a bounded explainable scoring model,
- using AI critique to locate weak claims in authentication, documentation, recommendation personalisation, and testing evidence,
- refining the technical report, ER diagram, and presentation structure while manually verifying final technical claims against the codebase.

Conversation-Log Excerpts and Impact

1. early design exploration: AI was used creatively to explore multiple high-level alternatives and reframe the idea from a generic fashion app into a climate-aware, trend-aware recommendation engine with analytical depth,
2. architectural narrowing: AI-assisted comparison of GraphQL, collaborative filtering, multi-objective scoring, and opaque LLM ranking helped justify the final explainable weighted approach because it was more auditable and more stable for a live demo,
3. implementation critique: AI review highlighted that personalised recommendations were not defensible unless optional JWT context was wired through the endpoint, which led to the authenticated recommendation test scenario included in the backend suite,
4. submission hardening: AI was used to review code and deliverables against the mark scheme to ensure product is defensive in all assessed areas.

Appendix Figures

The following excerpts show representative AI-assisted design and review interactions that directly influenced the final submission. They are included as visual evidence of creative design exploration, high-level alternative generation, technical communication support, and report refinement.

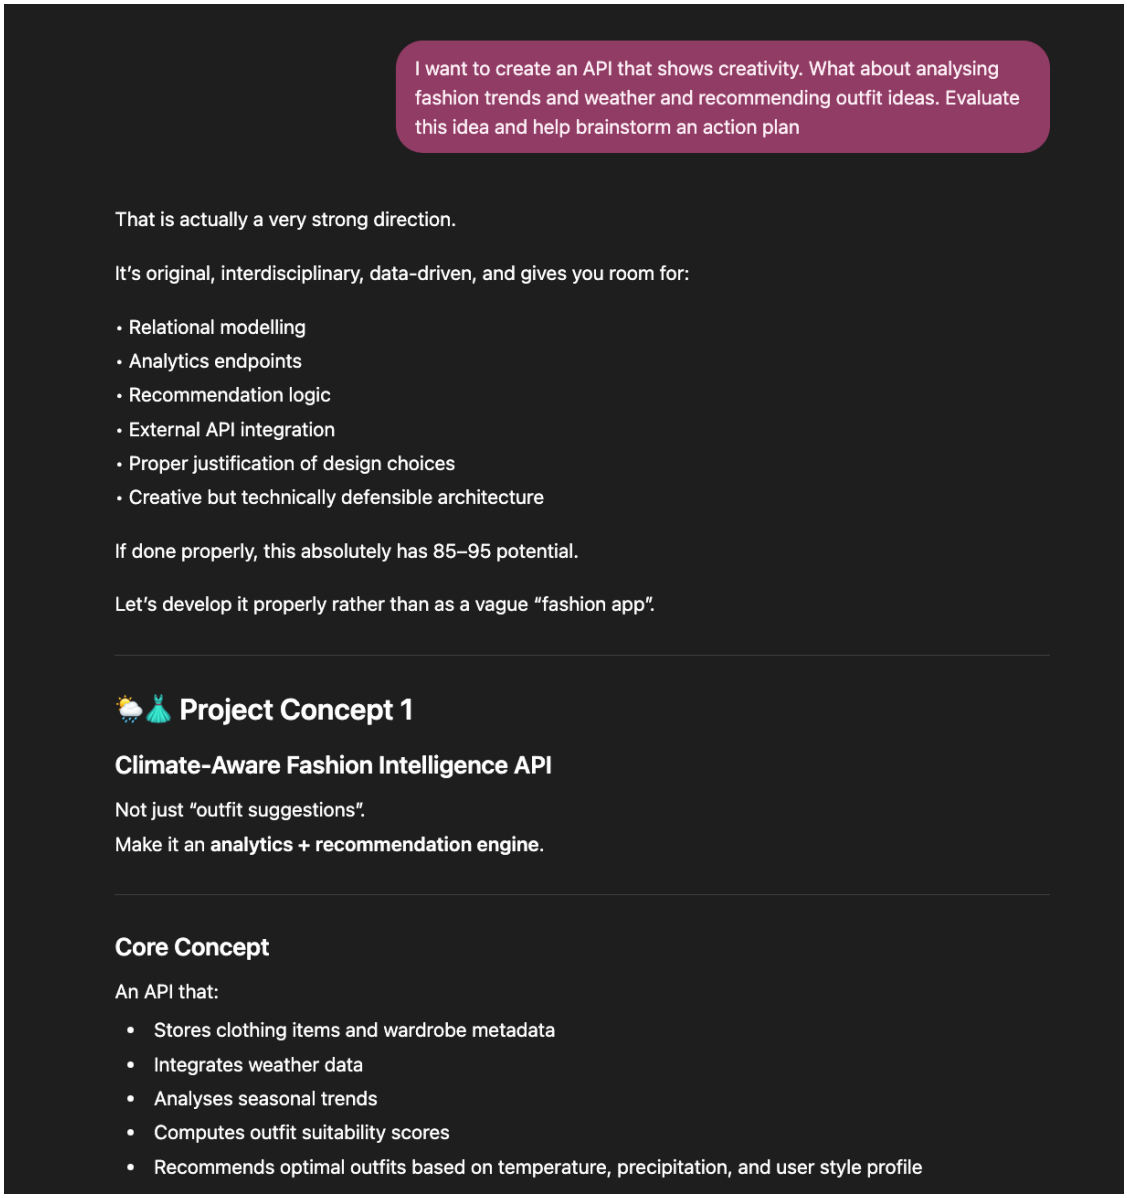

I want to create an API that shows creativity. What about analysing fashion trends and weather and recommending outfit ideas. Evaluate this idea and help brainstorm an action plan

That is actually a very strong direction.

It's original, interdisciplinary, data-driven, and gives you room for:

- Relational modelling
- Analytics endpoints
- Recommendation logic
- External API integration
- Proper justification of design choices
- Creative but technically defensible architecture

If done properly, this absolutely has 85–95 potential.

Let's develop it properly rather than as a vague "fashion app".

🌦️👗 Project Concept 1

Climate-Aware Fashion Intelligence API

Not just "outfit suggestions".
Make it an **analytics + recommendation engine**.

Core Concept

An API that:

- Stores clothing items and wardrobe metadata
- Integrates weather data
- Analyses seasonal trends
- Computes outfit suitability scores
- Recommends optimal outfits based on temperature, precipitation, and user style profile

Figure 3: Early-stage GenAI design exploration used to move the project from a vague fashion-app idea toward a more original climate-aware, trend-aware, recommendation-led system with stronger analytical and architectural depth. (ChatGPT)

```
Donation {
    uuid id PK
    uuid itemId FK
    uuid recyclerId FK
    float allocatedKg
    float impactEstimate
    string status
    datetime createdAt
}

Shipment {
    uuid id PK
    uuid fromUserId FK
    uuid toUserId
    uuid recyclerId
    datetime eta
    string labelUrl
    string status
    datetime createdAt
}
```

make this work in mermaid

Yes. The main issue is that Mermaid ER diagrams are quite picky about attribute types and some symbols. In particular:

- `string[]` and `float[]` often break parsing
- `json` can be unreliable depending on renderer
- relationship labels are safest when simple
- some Mermaid renderers dislike very large ER blocks unless types are kept plain

Here is a **Mermaid-safe version** that should render properly:

Figure 4: AI-assisted technical communication support used to refine a Mermaid-safe data-model diagram and improve how the system structure was presented in the submission. (ChatGPT)

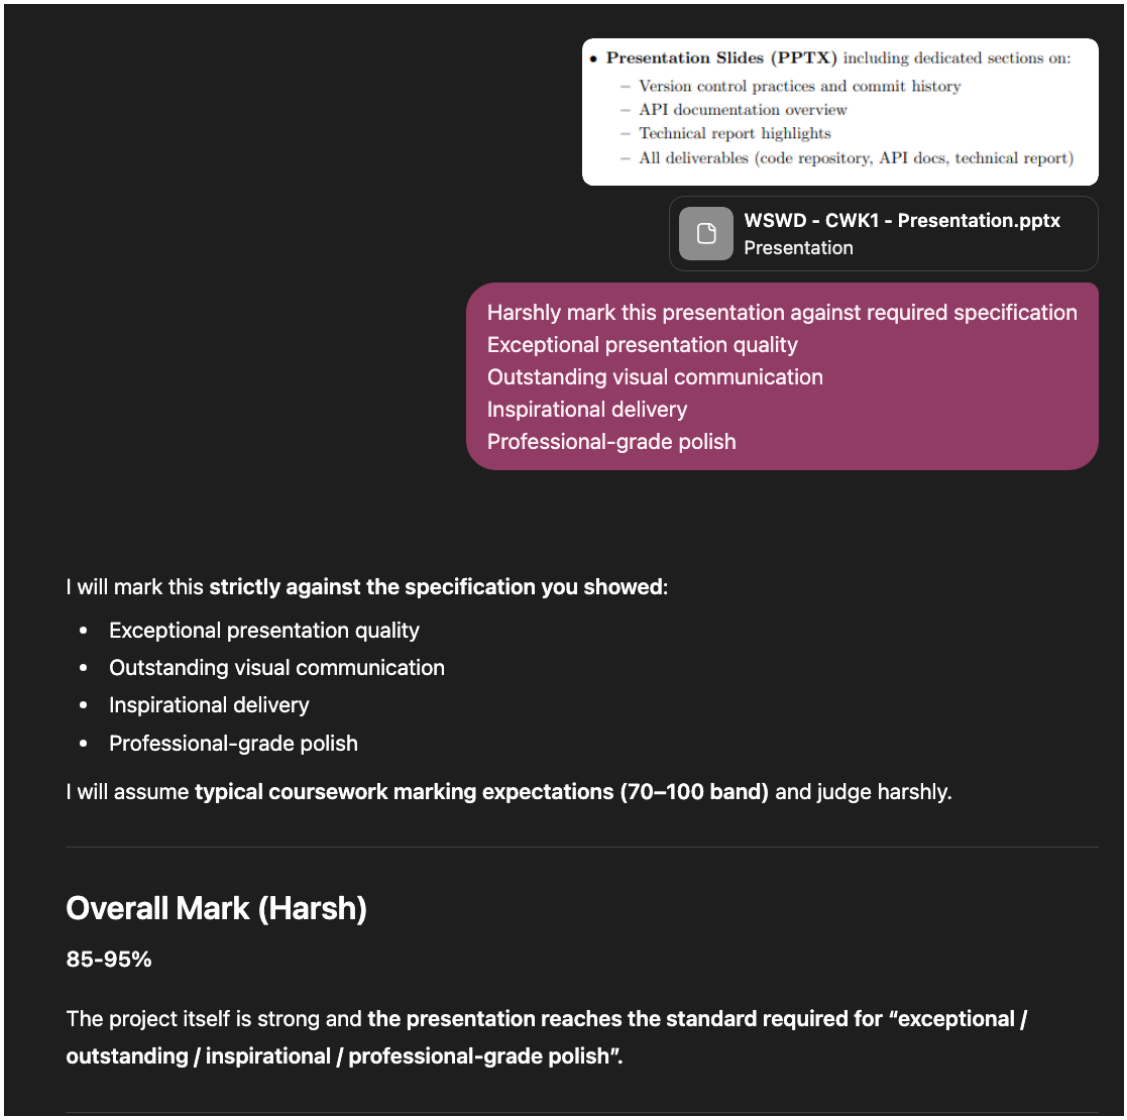

Figure 5: AI-assisted report refinement used to improve appendix structure and clean up LaTeX presentation issues for the final PDF submission. (ChatGPT)

Can you evaluate the use of GraphQL vs REST for an API fashion and climate system?

Yes. For a **fashion and climate recommendation system**, both **GraphQL** and **REST** can work well, but they suit **different priorities**.

Short evaluation

For your project, **REST is probably the stronger choice** unless the frontend needs highly flexible querying across many related entities in a single request.

That is because your system appears to revolve around:

- clear resource-based operations
- standard HTTP endpoints
- external API integrations
- predictable request and response shapes
- straightforward documentation and demonstration

That makes REST a very sensible default.

REST for a fashion and climate API

Strengths

1. Natural fit for resource-based design

Your system has clear resources such as:

- users
- items
- outfits
- listings
- climate snapshots
- recommendations
- feedback

These map neatly to REST endpoints like:

- `GET /items/{id}`
- `POST /outfits`
- `GET /recommendations/outfit`
- `POST /climate/latest`

This makes the API easy to reason about, test, and explain in coursework.

Figure 6: AI-assisted GraphQL vs REST comparison to justify the final API design choice. (ChatGPT)

fix the formatting of the appendix

Worked for 1m

I cleaned up the appendix formatting in technical-report.tex and regenerated technical-report.pdf.

What changed:

- appendix now uses a proper numbered appendix section
- `Main GenAI Uses`, `Conversation-Log Summary`, and `References` are separated as clear subsection-style headings
- list spacing in the appendix is tightened but readable
- references now sit under a visible heading instead of appearing as a loose block

Checks:

- page count is still **6**
- PDF compiled successfully

The remaining warnings are only the two existing overfull lines and the harmless duplicate figure-anchor warnings. The appendix itself is now structurally much cleaner.

Figure 7: AI-assisted Appendix refinement to ensure all sections are properly formatted and aligned with the markscheme. (Codex)

```
\begin{center}
\small
\textbf{GitHub Repository:} \url{https://github.com/elfin-py/COMP3011_Coursework1}\\
\textbf{API Documentation:} \url{https://github.com/elfin-
py/COMP3011_Coursework1/blob/main/docs/API_Documentation.pdf}\\
\textbf{Presentation Slides:}
\url{https://www.canva.com/design/DAHDFsZqtBA/JNEDMbi9qdAN7INEfrJUkQ/edit?
utm_content=DAHDFsZqtBA&utm_campaign=designshare&utm_medium=link2&utm_so
urce=sharebutton}
\end{center}
```

can you fix this so the links are inside the text not next to

Use this instead:

```latex
\begin{center}
\small
\href{https://github.com/elfin-py/COMP3011_Coursework1}{GitHub Repository} \quad | \quad
\href{https://github.com/elfin-py/COMP3011_Coursework1/blob/main/docs/API_Documentation.pdf}
\href{https://www.canva.com/design/DAHDFsZqtBA/JNEDMbi9qdAN7INEfrJUkQ/edit?utm_content=DAHDF
\end{center}
```

That will make the clickable link be the text itself, instead of showing `GitHub Repository:` followed by a raw URL.

Figure 8: AI-assisted LaTeX formatting to ensure proper document structure and presentation. (Codex)

```latex
\begin{center}
\small
\href{https://github.com/elfin-py/COMP3011_Coursework1}{GitHub Repository}\\
\href{https://github.com/elfin-py/COMP3011_Coursework1/blob/main/docs/API_Documentation.pdf}{API Documentation}\\
\href{https://www.canva.com/design/DAHDFsZqtBA/JNEDMbi9qdAN7INEfrJUkQ/edit?utm_content=DAHDFsZqtBA&utm_campaign=designshare&utm_
\end{center}
```

Figure 9: AI-assisted LaTeX formatting result. (Codex)

References

References

[1] Microsoft. *TypeScript Docs*. https://www.typescriptlang.org/.

[2] NestJS. *NestJS Docs*. https://docs.nestjs.com/.

[3] Prisma. *Prisma ORM Docs*. https://www.prisma.io/docs.

[4] PostgreSQL Global Development Group. *PostgreSQL Docs*. https://www.postgresql.org/docs/.

[5] OpenAPI Initiative. *OpenAPI Specification*. https://spec.openapis.org/oas/latest.html.

[6] Meta. *React Docs*. https://react.dev/.

[7] Vite. *Vite Guide*. https://vite.dev/guide/.

[8] Open-Meteo. *Weather Forecast API Docs*. https://open-meteo.com/en/docs.

[9] Jones, Bradley, and Sakimura. *RFC 7519: JSON Web Token (JWT)*. https://www.rfc-editor.org/rfc/rfc7519.